

# Problem Set 3

---

**Note** Your commit history on this assignment must show *incremental development* across at least two days. One single commit with all of your code is not acceptable.

## Contents


- [Python - Molecules as Graphs](#)
  - [Questions](#)
  - [Files](#)
- [C++ - Generic printing function](#)
  - [Overloading the stream insertion operator](#)

## Python - Molecules as Graphs

So far in this course, we've talked about running molecular mechanics and quantum chemistry simulations of molecules. But, how do you store information about molecules in databases? Or, if you have a set of molecules, how do you determine which molecules are similar to one another? When you have a large amount of data, visual inspection to determine similarity or equivalence is impractical. It also requires high skill and may be prone to errors. Thus, for cheminformatics or machine learning applications, we have to be able to express a molecule mathematically.

Many of the representations for small molecules that are used in cheminformatics and drug discovery are based on representing molecules using graph theory. There are algorithms and molecular representations that build on top of this concept.

A mathematical graph is made up of "nodes" or "vertices" (circles in the image below) and "edges" which connect nodes (line between circles in the image below).



**Figure 1** - A depiction of a graph with six nodes and seven edges. [source](#)

When molecules are represented as graphs, the atoms are represented as nodes in the graph, and the bonds are represented as edges.

In this homework, you will use a Python library called [NetworkX](#) for graph representation. You will also implement your own ring finding algorithm. However, if you were working on a real molecular application, you would likely use a specialized library like [RDKit](#) instead (we will use this library in a later lab!).

Your task for this homework is to use Python to represent molecular information read from a [Structured data File \(sdf\)](#) as a graph. You have been provided with a function called `parse_sdf` in `read.py` which will return a list of elements and bonds from an sdf file.

You will use a Python library called [NetworkX](#) to create a graph from molecular information read from an SDF file. You can then use graph functions to determine things about the molecular structure, such as the presence of rings.

Your code should create a graph using the information from the sdf file. You should then print out the number of rings in the molecule and the number of atoms in each ring. You should also use NetworkX to create a visualization of the network and save it to a file.

As an extra credit option, you will also implement your own code for counting the number of rings in a molecular graph using a depth-first search traversal. This extra credit is worth **10 points** on your Lecture Exercises (if you missed a week, this is a good way to make up points).

You might find it useful to work in a Jupyter notebook for this homework, but **your final code should be turned in as a .py file**.

Your code should include:

1. A function to create a NetworkX graph from the output of `parse_sdf`. Each atom should be a node and each bond should be an edge.
2. Calculate the number of rings (cycles) in the molecule and the number of atoms in each ring using NetworkX and a depth-first graph traversal. You will have to consult the NetworkX documentation to find an appropriate method. Have your script print the molecule name and number of rings. For each ring, print the number of atoms (nodes) in the ring.
3. Create a visualization of your molecule with `nx.draw_networkk`. Label the nodes in the visualization with the atom elements. You should also color the nodes by element using [CPK coloring](#).
4. **Extra Credit (10 Lecture Exercise points)** - Implementation of a ring finding algorithm (depth-first search) to find the number of rings in a molecule using your NetworkX graph. Note that the depth-first search approach will only work for simple molecules. You can see pseudocode for [a depth-first ring finding algorithm on Wikipedia](#). Use this as a guide to complete your task. Clearly label your extra credit task in a file called `extra_credit.py`.

## Questions

Answer these questions in your `README.md`.

1. What is an important feature of a NetworkX node? What data type did you choose to represent a node, and why? If you did not include information about the atom identity in your node, what type of Python data type could you have used to do that?
2. What is a depth-first search algorithm? Explain how the algorithm works (you may choose to use a combination of illustration and text). The depth-first search uses recursion - what is recursion and why is it used for this algorithm?

3. Use PubChem to get an SDF file for a molecule of choice and use your code to analyze it. What molecule did you choose and why?

## Files

Include the following files in your repo:

1. Your code which can create a NetworkX network from information in an SDF file.
2. A `README.md` which explains the repo purpose and how to run the code in your project.
3. A `Makefile` with the following targets:
  1. `environment` - creates the Python environment needed to run your code. Note that you will need to create an install libraries you need like NetworkX.
  2. `analyze` - analyze your molecule of choice.
  3. `clean` - remove images from `analyze`
  4. `lint` - run black and flake8

If you're interested in learning more about molecular representations, you might consider checking out the following review: ["Molecular representations in AI-driven drug discovery: a review and practical guide"](#)

## C++ - Generic printing function

Write a (templated) function that takes in an `std::vector` containing any type, and then loops over it and prints all the contents, with each element of the vector on its own line.

### Overloading the stream insertion operator

In C++, you can overload the stream insertion operator (`<<`). You can do this by writing a function that takes in two arguments, one being the output stream object, and the second argument being what is going to be inserted. In this case, the first argument is a generic `std::ostream` (output stream) object, and the second is what you want to print.

We will cover streams in a little more detail in Week 5. However, we will just say now that `std::cout` is a type of `std::ostream`. Therefore, by overloading this operator we can say `std::cout << vec`.

The function should return the `std::ostream` object that was passed in. This is what allows chaining calls to the operator `<<`. Of course, this function can be templated!

```
std::ostream & operator<<(std::ostream & os, ...)  
{  
    return os;  
}
```